

## Lecture 22

- ❖ Sections 5.1, 5.2, 5.2.1 and 5.2.2
- ❖ Overview of network layer
  - Forwarding and control planes
  - Network service models
- ❖ What is inside a router
  - Input port processing & destination-based forwarding
  - Switching

Network Layer Control Plane 4-1

1

## Chapter 5: Network Layer Control Plane

### Chapter objectives:

- understand principles behind network control plane:
  - traditional routing algorithms
  - network management, configuration
- instantiation, implementation in the Internet:
  - OSPF, DV (RIP), BGP
  - Internet Control Message Protocol: ICMP

Network Layer Control Plane 4-2

2

# Chapter 5: Network Layer

## 5.1 Introduction

Network Layer Control Plane 4-3

3

## Network-layer functions

- **forwarding**: move packets from router's input to appropriate router output *data plane*
- **routing**: determine route taken by packets from source to destination *control plane*


Two approaches to structuring network control plane:

- ❖ per-router control (traditional)
- ❖ logically centralized control (software defined networking)

4

## Per-router control plane

Individual routing algorithm components *in each and every router* interact in the control plane



5

## Chapter 5: Network Layer

### 5.1 Introduction

### 5.2 Routing algorithms


Network Layer Control Plane 4-6

6

## Routing protocols

**Routing protocol goal:** determine "good" paths (equivalently, routes), from sending hosts to receiving host, through network of routers

- ❖ **path:** sequence of routers packets traverse from given initial source host to final destination host
- ❖ **"good":** least "cost", "fastest", "least congested"
- ❖ **routing:** a "top-10" networking challenge!




7

## Routing protocols

**Routing protocol goal:** determine "good" paths (equivalently, routes), from sending hosts to receiving host, through network of routers

- ❖ **path:** sequence of routers packets traverse from given initial source host to final destination host
- ❖ **"good":** least "cost", "fastest", "least congested"
- ❖ **routing:** a "top-10" networking challenge!



8

## Routing Algorithm classification

### Global or decentralized information?

#### Global:

- ❖ all routers have complete topology, link cost info
- ❖ recursive computation
- ❖ "link state" algorithms

#### Decentralized:

- ❖ router knows physically-connected neighbors, link costs to neighbors
- ❖ iterative process of computation, exchange of info with neighbors
- ❖ "distance vector" algorithms

### Static or dynamic?

#### Static:

- ❖ routes change slowly over time


#### Dynamic:

- ❖ routes change more quickly
  - periodic update
  - in response to link cost changes
- ❖ used in Internet

Network Layer Control Plane 4-9

9

## Graph abstraction



Graph:  $G = (N, E)$


$N$  = set of routers =  $\{ u, v, w, x, y, z \}$

$E$  = set of links =  $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$

Network Layer Control Plane 4-10

10

## Graph abstraction: costs



•  $c(x, x')$  = cost of link  $(x, x')$

- e.g.,  $c(w, z) = 5$

• cost could always be 1, or  
inversely related to bandwidth,  
or inversely related to  
congestion

Cost of path  $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

Question: What's the least-cost path between u and z ?

Routing algorithm: algorithm that finds least-cost path

Network Layer Control Plane 4-11

11

## Chapter 5: Network Layer

5.1 Introduction

5.2 Routing algorithms

5.2.1 link-state routing algorithm

Network Layer Control Plane 4-12

12

## A Link-State Routing Algorithm

### Dijkstra's algorithm

- ❖ net topology, link costs known to all nodes
  - accomplished via "link state broadcast"
  - all nodes have same info
- ❖ computes least cost paths from one node ("source") to all other nodes
  - gives *forwarding table* for that node
- ❖ iterative: after k iterations, know least cost path to k dest.'s

### Notation:

- ❖  $c(x,y)$ : link cost from node x to y;  $= \infty$  if not direct neighbors
- ❖  $D(v)$ : current value of cost of path from source to dest. v
- ❖  $p(v)$ : predecessor node along path from source to v
- ❖  $N'$ : set of nodes whose least cost path definitively known

Network Layer Control Plane 4-13

13

## Overview of Dijkstra's Algorithm:

- ❖ The algorithm starts by calculating the known direct distances from the source,  $u$ , to all nodes,  $v$ ,  $D(v)$ , which is equal to  $c(u,v)$ .
- ❖ The algorithm selects the closest node to  $u$ ,  $v$  (with minimum  $D(v)$ ) and adds it to  $N'$
- ❖ The algorithm recalculates the minimum distance to the nodes not in  $N'$  as the minimum of the already known distance and the distance through the newly added node, & updates the routes accordingly

Network Layer Control Plane 4-14

14

## Dijkstra's Algorithm

```

1 Initialization:
2  N' = {u}
3  for all nodes v
4    if v adjacent to u
5      then D(v) = c(u,v)
6    else D(v) = ∞
7
8 Loop
9  find w not in N' such that D(w) is a minimum
10 add w to N'
11 update D(v) for all v adjacent to w and not in N' :
12   D(v) = min( D(v), D(w) + c(w,v) )
13 /* new cost to v is either old cost to v or known
14    shortest path cost to w plus cost from w to v */
15 until all nodes are in N'

```

Network Layer Control Plane 4-15


15

## Dijkstra's algorithm: example

Step	N'	D(v) p(v)	D(w) p(w)	D(x) p(x)	D(y) p(y)	D(z) p(z)
0	u	7,u	3,u	5,u	∞	∞
1	uw	6,w		5,u	11,w	∞
2	uwx	6,w			11,w	14,x
3	uwxv				10,v	14,x
4	uwxvy					12,y
5	uwxvyz					

### Notes:

- ❖ construct shortest path tree by tracing predecessor nodes
- ❖ ties can exist (can be broken arbitrarily)




Network Layer Control Plane 4-16

16



## Dijkstra's algorithm: another example

Step	N'	D(v),p(v)	D(w),p(w)	D(x),p(x)	D(y),p(y)	D(z),p(z)
0	u	2,u	5,u	1,u	$\infty$	$\infty$
1	ux	2,u	4,x		2,x	$\infty$
2	uxy	2,u	3,y			4,y
3	uxyv		3,y			4,y
4	uxyvw					4,y
5	uxyvwz					




Network Layer Control Plane 4-17

17

## Dijkstra's algorithm: example (2)

Resulting shortest-path tree from u:



Resulting forwarding table in u:

destination	link
v	(u,v)
x	(u,x)
y	(u,x)
w	(u,x)
z	(u,x)

Network Layer Control Plane 4-18

18

## Dijkstra's algorithm, discussion

Algorithm complexity:  $n$  nodes

- ❖ each iteration: need to check all nodes,  $w$ , not in  $N$
- ❖  $n(n+1)/2$  comparisons:  $O(n^2)$
- ❖ more efficient implementations possible:  $O(n \log n)$

Network Layer Control Plane 4-19

19

## Chapter 5: Network Layer

5.1 Introduction

5.2 Routing algorithms

5.2.1 link-state routing algorithm

5.2.2 distance vector routing

Network Layer Control Plane 4-20

20

## Distance Vector Algorithm

### Bellman-Ford Equation (dynamic programming)

Define

$d_x(y) :=$  cost of least-cost path from  $x$  to  $y$

Then


$$d_x(y) = \min_v \{c(x,v) + d_v(y)\}$$

where min is taken over all neighbors  $v$  of  $x$

Network Layer Control Plane 4-21

21

## Bellman-Ford example



Source is  $u$ , and wants to reach  $z$   
Neighbors of  $u$  are  $v$ ,  $w$  and  $x$

Clearly,  $d_v(z) = 5$ ,  $d_x(z) = 3$ ,  $d_w(z) = 3$

B-F equation says:

$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$

Node that achieves minimum is next  
hop in shortest path (node  $x$ )  $\rightarrow$  forwarding table

Network Layer Control Plane 4-22

22

## Distance Vector Algorithm

- ❖  $D_x(y)$  = estimate of least cost from  $x$  to  $y$ 
  - $x$  maintains distance vector  $D_x = [D_x(y): y \in N]$
- ❖ node  $x$ :
  - knows cost to each neighbor  $v$ :  $c(x,v)$
  - maintains its neighbors' distance vectors.  
For each neighbor  $v$ ,  $x$  maintains  $D_v = [D_v(y): y \in N]$

Network Layer Control Plane 4-23

23

## Distance vector algorithm (4)

### Basic idea:

- ❖ from time-to-time, each node sends its own distance vector estimate to neighbors
- ❖ when  $x$  receives new DV estimate from neighbor, it updates its own DV using B-F equation:
$$D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\} \quad \text{for each node } y \in N$$
- ❖ under minor, natural conditions, the estimate  $D_x(y)$  converge to the actual least cost  $d_x(y)$

Network Layer Control Plane 4-24

24

## Distance Vector Algorithm (5)

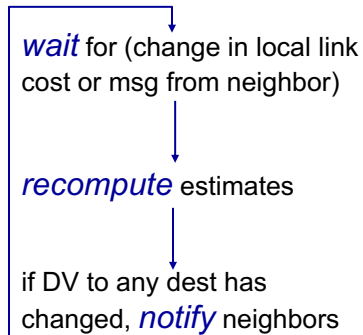
### Distributed:

- ❖ each node notifies neighbors *only* when its DV changes
  - neighbors then notify their neighbors if necessary

### Iterative, asynchronous:

- each local iteration caused by:
  - ❖ local link cost change
  - ❖ DV update message from neighbor

### Each node:




Network Layer Control Plane 4-25

25

## Distance Vector Routing: Example


(X,i):  
X: next hop to N1  
i: distance to N1



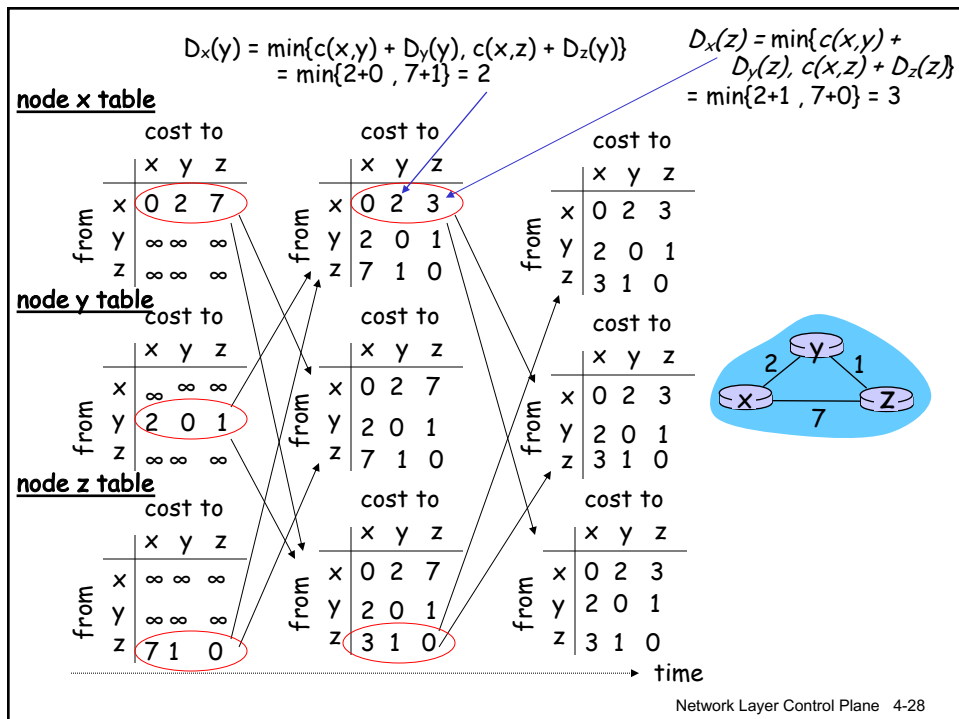
Iter.	A	B	C	D	E	F
0	(-,1)	(-,4)	(•,∞)	(-,8)	(-,1)	(•,∞)
1	(-,1)	(A,3)	(B,9)	(E,6)	(-,1)	(D,28)
2	(-,1)	(A,3)	(D,7)	(E,6)	(-,1)	(C,10)
3	(-,1)	(A,3)	(D,7)	(E,6)	(-,1)	(C,8)

Network Layer Control Plane 4-26

26



27




28

## Distance Vector: link cost changes

### Link cost changes:

- ❖ node detects local link cost change
- ❖ updates routing info, recalculates distance vector
- ❖ if DV changes, notify neighbors



"good  
news  
travels  
fast"

$t_0$ : y detects link-cost change, updates its DV, informs its neighbors.

$t_1$ : z receives update from y, updates its table, computes new least cost to x, sends its neighbors its DV.

$t_2$ : y receives z's update, updates its distance table. y's least costs do *not* change, so y does *not* send a message to z.


Network Layer Control Plane 4-29

29

## Distance Vector: link cost changes


### Link cost changes:

- ❖ **bad news travels slow** - "count to infinity" problem!
- ❖ Many iterations before algorithm stabilizes: how many?



### Poisoned reverse:

- ❖ If Y routes through X to get to W :
  - Y tells X its (Y's) distance to W is infinite (so X won't route to W via Y)
- ❖ will this completely solve count to infinity problem?



Network Layer Control Plane 4-30

30

## Comparison of LS and DV algorithms

### Message complexity

- ❖ LS: with  $n$  nodes,  $E$  links,  $O(nE)$  msgs sent
- ❖ DV: exchange between neighbors only
  - convergence time varies

### Speed of Convergence

- ❖ LS:  $O(n^2)$  algorithm requires  $O(nE)$  msgs
  - may have oscillations
- ❖ DV: convergence time varies
  - may be routing loops
  - count-to-infinity problem

### Robustness: what happens if router malfunctions?

#### LS:

- node can advertise incorrect *link* cost
- each node computes only its *own* table

#### DV:

- DV node can advertise incorrect *path* cost
- each node's table used by others
  - error propagate thru network

Network Layer Control Plane 4-31